

1 This document provides a design description of the FORTH VM support being
2 included in the M65C02A soft-core processor. The M65C02A soft-core processor
3 is an extended version of the 6502/65C02 microprocessor originally developed
4 by Western Design Center (WDC).
5
6 The purpose of this document is to document the design decisions made with
7 respect to custom instructions added to the basic M65C02A instruction set in
8 order to provide better support for a FORTH VM than a standard 65C02-
9 compatible processor. It is an objective of this effort to add the least
10 number of instructions (and dedicated hardware) to the M65C02A instruction set
11 as necessary to provide an efficient FORTH VM.
12
13 The instruction set of the M65C02A is already extended beyond that of a
14 standard 6502/65C02 microprocessor. The M65C02A implements all of the standard
15 6502/65C02 instructions and addressing modes, but it also includes the four
16 bit-oriented instructions added by Rockwell and the WAI and STP instructions
17 included by WDC in its extended 6502/65C02 implementation: the W65C02S.
18
19 The M65C02A adds 8 base pointer-relative and 8 post-indexed base pointer-
20 relative indirect instructions matching the W65C816 processor's 16 stack-
21 relative/stack-relative indirect instructions. The M65C02A also includes the
22 W65C816's three 16-bit push instructions and its 16-bit relative branch
23 instruction. But the M65C02A adds an additional 16-bit push instruction,
24 several 16-bit pop/pull instructions, post-indexed base pointer-relative
25 indirect jumps and subroutine calls, and a 16-bit relative subroutine call
26 instruction. The M65C02A also includes an instruction for supporting co-
27 processors, and six prefix instructions that change the operand size, add
28 indirection, and override the destination registers of various instructions.
29
30 The IND prefix instruction provides the programmer with the ability to add
31 indirection to many instructions which do not include an indirect addressing
32 mode. The SIZ and the OAX/OAX prefix instructions provide the programmer the
33 means to increase the operand/ALU operation size, and to override the default
34 destination register of an instruction. The OSX prefix instruction allows the
35 X register to be used as an auxiliary hardware stack pointer in memory. The
36 effects of IND and SIZ are combined in the ISZ prefix instruction.
37
38 The FORTH VM can generally be thought of as being constructed from a minimum
39 of two stacks: (1) a parameter/data stack (PS), and (2) a return stack (RS).
40 The PS is intended to hold all parameters/data used within a program/function,
41 and the RS generally holds the return addresses of FORTH program words. In
42 addition to holding FORTH program word return addresses, the RS is also used
43 to hold loop addresses, and may be used to hold parameter/data addresses.
44
45 These stacks are generally implemented within the memory of whatever
46 microprocessor is implementing the FORTH VM. To support an efficient
47 implementation of the FORTH VM, any specific FORTH implementation should
48 provide hardware assisted stack pointers whenever possible.
49
50 The M65C02A soft-core processor is an implementation that provides a faithful
51 reimplementaion of the 6502/65C02 processor in a synthesizable manner.
52 Although it provides all of the registers and executes all of the standard
53 instructions, its implementation attempts to provide the best performance
54 possible. As a result, the M65C02A reimplementaion is not instruction cycle
55 length compatible with 6502/65C02 or W65C02S processors.
56
57 In all other respects it is a 6502/65C02 microprocessor. That means that, in
58 its 6502/65C02 compatibility mode, it appears to the programmer as having an
59 8-bit arithmetic accumulator (A), an 8-bit processor flags registers (P), two
60 8-bit index registers (X, Y), and a 8-bit system stack pointer (S). Except for
61 S, the general purpose accumulator (A) and the index registers (X, Y) do not
62 provide any hardware support for the FORTH VM data or return stacks. In its
63 extended mode, the M65C02A, has three modified 16-bit push-down stacks, one
64 for each base register: A, X, Y. Each register stack is composed of three 16-
65 bit registers arranged such that the top-of-stack is the register directly
66 affected by the load/store and ALU instructions. Automatic pushing/popping of
67 each register stack is not performed. Instead, explicit management of each
68 register stack is left to the programmer. Three specific instructions are
69 available for managing the register stacks: DUP (duplicate/push TOS), SWP
70 (swap TOS and NOS), ROT (rotate the register stack). Thus, to load the TOS
71 register and to push down the values currently in the stack, a DUP instruction
72 must precede a load instruction. To simply duplicate the TOS and push down the
73 stack, only a DUP instruction is needed. To store the TOS value to memory and
74 pop the stack, a store instruction followed by a ROT instruction is needed.
75
76 Although the extensions to the 6502/65C02 instruction set provided by the
77 M65C02A are uniformly applicable to most instructions and registers, each
78 register stack have some unique capabilities. For example, the A register
79 stack supports byte swapping and wrap-around nibble rotation (4 bits at a

80 time) in the stack's TOS register. Similarly, the TOS of the X register stack
 81 can function as a third hardware-assisted stack pointer and as a automatically
 82 incremented or decremented memory pointer. Finally, the TOS of the Y register
 83 stack may function as an automatically incremented or decremented memory
 84 pointer.
 85

86 Performance degradations in a 6502/65C02 FORTH VM are generally due to the
 87 fact that all of the registers are 8-bits in length and 16 bits is the assumed
 88 operand and pointer size of the VM. Support is provided in the 6502/65C02
 89 instruction set architecture for multi-precision addition and subtraction, but
 90 any operations greater than 8 bits in length will entail several loads and
 91 stores. All of the additional steps necessary to implement a 16-bit or 32-bit
 92 FORTH VM operations reduces the performance a native 6502/65C02 FORTH VM can
 93 deliver.
 94

95 Brad Rodriguez wrote a series of articles on the development of FORTH VMs. The
 96 lead article of the series, "MOVING FORTH Part 1: Design Decisions in the
 97 Forth Kernel" (<http://www.bradrodriguez.com/papers/moving1.htm>), provides a
 98 good discussion of the design trades needed when implementing a FORTH VM on a
 99 microprocessor.
 100

101 Brad Rodriguez identifies the following registers as being the "classic" FORTH
 102 VM registers:
 103

104	W	- Work Register	: general work register
105	IP	- Interpreter Pointer	: address of the next FORTH word to execute
106	PSP	- Parameter Stack Pointer	: points to the top of parameter/data stack
107	RSP	- Return Stack Pointer	: points to the return address
108	UP	- User Pointer	: points to User space of a multi-task FORTH
109	X	- eXtra register	: temporary register for next address

110

111 There are several generally accepted methods for implementing the FORTH VM.
 112 The classic FORTH VM is implemented using a technique known as Indirect
 113 Threaded Code (ITC). The next most common approach is a technique known as
 114 Direct Threaded Code (DTC). A more recent approach is a technique known as
 115 Subroutine Threaded Code (STC). A final, less commonly used approach is a
 116 technique known as Token Threaded Code (TTC).
 117

118 A TTC FORTH VM uses tokens that are smaller than the basic address pointer
 119 size to refer to FORTH words. The smaller size of the tokens allows a FORTH
 120 program to be compressed into a smaller image. The additional indirection
 121 required to locate the memory address of the FORTH word referenced by a
 122 particular token makes TTC FORTH VM implementations the slowest of the FORTH
 123 VM techniques. The pre-indexed indirect addressing modes of the 6502/65C02 can
 124 be used to implement a token threaded VM using the M65C02A. Thus, the basic
 125 indexed addressing modes of the M65C02A provide the necessary support to
 126 implement a TTC FORTH VM.
 127

128 For an STC FORTH VM, each word is called using a native processor call
 129 instruction. Thus, there is no need for an IP register, and there is no inner
 130 interpreter. The FORTH program simply chains together the various FORTH words
 131 using native processor calls. There are two penalties for this simplicity: (1)
 132 subroutine calls are generally larger than simple address pointers; and (2)
 133 requires pushing and popping the return stack on entry to and exit from each
 134 FORTH word. Thus, STC FORTH programs may be larger, and the additional
 135 push/pop operations performed by the native subroutine calling instructions
 136 may not deliver the performance improvements expected. Since an STC FORTH VM
 137 relies on the basic instructions of the processor, and the M65C02A provides
 138 those instructions, no additional instructions are needed in the M65C02A's
 139 instruction set to support an STC FORTH.
 140

141 ITC FORTH and DTC FORTH VMs both require an inner interpreter. Thus, if the
 142 instruction set of a processor allows the implementation of the inner
 143 interpreter with a minimum number of instructions, then a FORTH VM on that
 144 processor would be "faster" than a FORTH VM on a processor without that
 145 support. This is the prime motivating factor for adding custom instructions to
 146 the M65C02A to support FORTH VMs.
 147

148 A TTC/DTC/ITC FORTH VM is composed of two interpreters: (1) an outer
 149 interpreter, and (2) an inner interpreter. The outer interpreter is written in
 150 FORTH, i.e. it is composed of FORTH words. The inner interpreter, on the other
 151 hand, "executes" FORTH words. Therefore, a TTC/DTC/ITC FORTH VM spends the
 152 majority of its processing time in its inner interpreter. Thus, any decrease
 153 in the number of clock cycles required to "execute" a FORTH word will result
 154 in a clear increase in the performance of a FORTH program all other things
 155 being equal.
 156

157 The basic structure of a 16-bit FORTH word is provided by the following C-like
 158 structure:

```

159
160     struct {
161         uint8_t      Len;
162         uint8_t      [Max_Name_Len] Name;
163         uint16_t     *Link;
164         uint16_t     *Code_Fld;
165         uint8_t      [Code_Len] Param_Fld;
166     } FORTH_word_t;
167

```

168 There are other forms, but the preceding structure defines all of the
 169 necessary components of the FORTH word, and succinctly conveys the required
 170 elements of a FORTH word.

171
 172 The first three fields provide the "dictionary" header of FORTH words in a
 173 FORTH program. The first field defines the length of the name of the FORTH
 174 word. This is used to distinguish two or more FORTH words in a FORTH program
 175 whose names share the same initial letters but which differ in length. The
 176 second field defines the significant elements of the name of the FORTH word.
 177 The total lengths of these two fields will determine the amount of memory that
 178 is required just for the dictionary of a FORTH program. These two fields are
 179 generally limited in size in order to conserve memory. If the immediate mode
 180 of the FORTH compiler is not supported or included in the distribution of a
 181 FORTH program, then the fields supporting the "dictionary" can be removed to
 182 recover their memory for use by the application.

183
 184 The fields, Code_Fld and Param_Fld, represent the "executable" portion of a
 185 FORTH word. There are two types of FORTH words: (1) secondaries, and (2)
 186 primitives. Secondaries are the predominant type of words in a FORTH program.
 187 Their Param_Fld doesn't contain any native machine code. The Param_Fld of
 188 FORTH secondaries is simply a list of the Code_Fld of other FORTH words.
 189 Primitives are the FORTH words that perform the actual work of any FORTH
 190 program. The Code_Fld of a primitive is a link to their Param_Fld, which
 191 contains the machine code that performs the work the primitive FORTH word is
 192 expected to provide.

193
 194 In FORTH, the outer interpreter is used to perform immediate operations, and
 195 construct, define, or compile other FORTH words. As already stated above, the
 196 FORTH VM's outer interpreter is generally composed of secondary FORTH words.
 197 After all of the FORTH words associated with a FORTH program have been
 198 compiled, the outer interpreter simply transfers control to the inner
 199 interpreter to "execute" the top-most FORTH word of the program.

200
 201 The FORTH VM's inner interpreter "executes" FORTH words. Since the outer
 202 interpreter is mostly composed of secondary FORTH words, the inner interpreter
 203 must move through each word until a FORTH primitive is found, and then
 204 transfer temporary control to the machine code. The machine code of the
 205 primitive FORTH word must return control to the inner interpreter once it
 206 completes its task.

207
 208 The inner interpreter "executes" the FORTH word that its IP points to. It must
 209 advance the IP through the Param_Fld of a secondary FORTH word, and jump to
 210 the machine code pointed to by the Code_Fld of a FORTH primitive. The Code_Fld
 211 of a secondary does not point to the Param_Fld of the word. Instead it points
 212 to an inner interpreter function that "enters" the Param_Fld. The Code_Fld of
 213 a primitive does point to the Param_Fld, and the inner interpreter simply
 214 jumps to the machine code stored in the Param_Fld of the word.

215
 216 Thus, there are three fundamental operations that the inner interpreter of a
 217 DTC/ITC FORTH VM must perform:

- 218 (1) NEXT : fetch the FORTH word addressed by IP; advance IP.
- 219 (2) ENTER : save IP; load IP with the Code_Fld value; perform NEXT.

220
 221 Each primitive FORTH word must transfer control back to the inner interpreter
 222 so that the next FORTH word can be "executed". This action may be described as:

- 223 (3) EXIT : restore IP; perform NEXT.

224
 225
 226 These three fundamental operations are very similar to the operations that the
 227 target processor performs in executing its machine code. NEXT corresponds
 228 directly to the normal instruction fetch/execute cycle of any processor. ENTER
 229 corresponds directly to the subroutine call of any processor. Similarly, EXIT
 230 corresponds directly to the subroutine return of any processor.

231
 232 As previously discussed, FORTH uses two stacks: parameter/data stack and
 233 return stack. Thus, ENTER and EXIT save and restore the IP from the return
 234 stack. Most processors implement a single hardware stack into which return
 235 addresses and data are placed. FORTH maintains strict separation between the
 236 parameter/data stack and the return stack because it uses a stack-based

238 arithmetic architecture. Mixing return addresses and parameters/data on a
 239 single stack would complicate the passing of parameters and their processing.
 240
 241 Table 6.3.1 in Koopman's "Stack Computers", provides a summary the relative
 242 frequency of the most frequently used FORTH words for several FORTH
 243 applications:
 244

245 Name	FRAC	LIFE	MATH	COMPILE	AVE
246 CALL	11.16%	12.73%	12.59%	12.36%	12.21%
247 EXIT	11.07%	12.72%	12.55%	10.60%	11.74%
248 VARIABLE	7.63%	10.30%	2.26%	1.65%	5.46%
249 @	7.49%	2.05%	0.96%	11.09%	5.40%
250 OBRANCH	3.39%	6.38%	3.23%	6.11%	4.78%
251 LIT	3.94%	5.22%	4.92%	4.09%	4.54%
252 +	3.41%	10.45%	0.60%	2.26%	4.18%
253 SWAP	4.43%	2.99%	7.00%	1.17%	3.90%
254 R>	2.05%	0.00%	11.28%	2.23%	3.89%
255 >R	2.05%	0.00%	11.28%	2.16%	3.87%
256 CONSTANT	3.92%	3.50%	2.78%	4.50%	3.68%
257 DUP	4.08%	0.45%	1.88%	5.78%	3.05%
258 ROT	4.05%	0.00%	4.61%	0.48%	2.29%
259 USER	0.07%	0.00%	0.06%	8.59%	2.18%
260 C@	0.00%	7.52%	0.01%	0.36%	1.97%
261 I	0.58%	6.66%	0.01%	0.23%	1.87%
262 =	0.33%	4.48%	0.01%	1.87%	1.67%
263 AND	0.17%	3.12%	3.14%	0.04%	1.61%
264 BRANCH	1.61%	1.57%	0.72%	2.26%	1.54%
265 EXECUTE	0.14%	0.00%	0.02%	2.45%	0.65%

266
 267 In table above, CALL corresponds to ENTER. As can be seen, the remaining
 268 common FORTH words are a combination of parameter stack operations (DUP, ROT,
 269 SWAP), parameter stack loads (VARIABLE, LIT, CONSTANT, @, C@), parameter stack
 270 arithmetic and logic operations (+, =, AND), parameter stack branching and
 271 looping (OBRANCH, BRANCH, I), and parameter and return stack operations (R>,
 272 >R), and special operations (USER, EXECUTE).
 273

274 With the previous discussion and the FORTH word frequency data in the
 275 preceding table it is easy to assert that the M65C02A instruction set should
 276 contain custom instructions for at least NEXT, ENTER, and EXIT. The question
 277 then becomes to what extent should these operations be supported? In other
 278 words, should they be supported by a single instruction each and should they
 279 be supported for both ITC and DTC FORTH VMs?
 280

281 The following pseudo code defines the operations for these three operations in
 282 terms of the ITC and the DTC models:
 283

284	ITC	DTC
285	=====	=====
286	NEXT: W <= (IP++) -- Ld *Code_Fld ;	W <= (IP++) -- Ld *Code_Fld
287	PC <= (W) -- Jump Indirect ;	PC <= W -- Jump Direct
288	=====	=====
289	ENTER: (RSP--) <= IP -- Push IP on RS ;	(RSP--) <= IP -- Push IP on RS
290	IP <= W + 2 -- => Param_Fld ;	IP <= W + 2 -- => Param_Fld
291	;NEXT	
292	W <= (IP++) -- Ld *Code_Fld ;	W <= (IP++) -- Ld *Code_Fld
293	PC <= (W) -- Jump Indirect ;	PC <= W -- Jump Direct
294	=====	=====
295	EXIT:	
296	IP <= (++RSP) -- Pop IP frm RS ;	IP <= (++RSP) -- Pop IP frm RS
297	;NEXT	
298	W <= (IP++) -- Ld *Code_Fld ;	W <= (IP++) -- Ld *Code_Fld
299	PC <= (W) -- Jump Indirect ;	PC <= W -- Jump Direct
300	=====	=====

301
 302 Except for the indirection needed for ITC, there are several key takeaways
 303 from the side-by-side comparison provided above of the NEXT, ENTER, and EXIT
 304 operations. First, ENTER and EXIT are essentially the same for ITC and DTC
 305 FORTH VMs. Second, both ENTER and EXIT terminate with the operations
 306 implemented by NEXT. Also note that EXIT is simply an RS pop operation followed
 307 by a DTC/ITC NEXT operation.
 308

309 Finally, there are two important observations made by Rodriguez:
 310

- 311 (1) if W is left pointing to the Code_Fld of the word being executed, the
 312 Param_Fld of a FORTH word being ENTERed can be found using the value in W;
 313
- 314 (2) providing a second stack pointer for the RS is important and will greatly
 315 improve the performance of the inner interpreter.
 316

317 The preceding analysis and discussions set the stage for the critical design
 318 decisions for an M65C02A FORTH VM:
 319
 320 (1) mapping the FORTH VM registers onto the M65C02A registers;
 321
 322 (2) determining how to modify the M65C02A to support the FORTH VM inner
 323 interpreter operations.
 324
 325 The IP register is strictly used as the instruction pointer of the inner
 326 interpreter. It cannot be assigned to the target processor's program counter,
 327 but it does operate as such for the inner interpreter. An easy means for
 328 including IP is to place it in zero page memory, but this means that several
 329 memory cycles will be needed to increment, push, pop, or otherwise manipulate
 330 its value. A better solution is to add a 16-bit register within the core.
 331 Alternatively, IP may be accessed with whatever custom instructions are added
 332 to the M65C02A instruction set to support NEXT, ENTER, and EXIT. In addition
 333 to load and store operations, support should be provided to increment the IP
 334 by 2.
 335
 336 Similarly, the W register is used strictly as a pointer for indirect access to
 337 a FORTH word by the inner interpreter. It is only loaded indirectly from IP.
 338 Like the IP register, it can easily be implemented in zero page memory, but
 339 this means that several memory cycles will be needed to increment, push, pop,
 340 or otherwise manipulate its value. Like the IP, the best way for the M65C02A
 341 to support W is to include it in the processor core itself. Also, for it to be
 342 effectively utilized, support should be provided to increment the W register
 343 by 2.
 344
 345 The PS is used more often than the RS. Thus, it makes more sense, from a speed
 346 perspective, to use the native M65C02A stack for the PS. Using a pre-indexed
 347 zero page location for the RSP will slow the push and pop operations
 348 significantly. Therefore, a better solution would be to use one of the index
 349 registers as the RSP, and place the RS anywhere in memory, including page 0.
 350
 351 As the 6502/65C02 index register for the less frequently used pre-indexed
 352 (direct and indirect) addressing modes, the X index register is the natural
 353 choice to provide a hardware-assisted RSP. The auxiliary stack pointer
 354 capabilities of the TOS of the X register stack easily allow X to be used as
 355 the RSP. In addition, two instructions will be added to support pushing and
 356 pulling the IP from either stack, and a single cycle instruction will be added
 357 to increment the IP by 1. By overloading the IND prefix instruction, it is
 358 possible to use the dedicated IP push, pop, and increment instructions to
 359 perform the same operations with the W register.
 360
 361 (Note: instead incurring a byte/cycle penalty by using the OSX prefix
 362 instruction be used before any FORTH VM instructions that use the RS, the OSX
 363 prefix instruction's effects have been redefined to override of the default
 364 stack pointer of any instruction that utilizes the stack. The default stack
 365 for the FORTH VM ENT, PHI, PLI, PHW, and PLW has been set to use the auxiliary
 366 stack provided by X. The change in the behavior of OSX means that only when
 367 the PS is needed will these five FORTH VM instruction require OSX. It also
 368 saves 1 byte/cycle for every access to the RS.)
 369
 370 Thus, the FORTH VM supported by the M65C02A will provide the following mapping
 371 of the various FORTH VM registers:
 372
 373 IP - Internal dedicated 16-bit register
 374 W - Internal dedicated 16-bit register
 375 PSP - System Stack Pointer (S), allocated in memory (page 1 an option)
 376 RSP - Auxiliary Stack Pointer (X), allocated in memory (page 0 an option)
 377 UP - Memory (page 0 an option)
 378 X - Not needed, {OP2, OP1} or MAR can provide temporary storage required
 379
 380 The inner interpreter of the FORTH VM will be implemented directly in the
 381 M65C02A using five dedicated instructions: NXT (NEXT), ENT (ENTER), PLI (Pull
 382 IP), PHI (Push IP), and INI (Increment IP by 1). All five of these
 383 instructions can be prefixed by IND. The NXT and ENT instructions directly
 384 support a DTC FORTH VM, and when prefixed by IND, they support an ITC FORTH
 385 VM. The DTC EXIT will be implemented using the PLI NXT instruction sequence,
 386 and the ITC EXIT will be implemented using the PLI IND NXT instruction
 387 sequence. Access and control of the IP is provided by the PHI, PLI, and INI
 388 instructions. (When these instructions are prefixed by IND, access and control
 389 of W is provided: PHW, PLW, and INW, respectively.)
 390
 391 Loading constants/literals is a frequent operation in FORTH programs. Thus,
 392 support for efficient loading of in-line constants/literals relative to the IP
 393 is included in the M65C02A. The LDA ip,I++ instruction will load the byte
 394 which follows the current IP into the accumulator and advances the IP by 1. If
 395 this instruction is prefixed by SIZ, then the word following the current IP is

396 loaded into the accumulator and the IP is advanced by 2. If prefixed by IND,
 397 the instruction becomes LDA (ip,I++), which uses the 16-bit word following the
 398 current IP as a byte pointer. The IP is advanced by 2, and the byte pointed to
 399 by the pointer is loaded into the accumulator. If prefixed by ISZ, the word
 400 following the current IP is used as a word pointer, while the IP is advanced
 401 by 2, to load a word into the accumulator.

402
 403 The LDA ip,I++ instruction is matched by the STA ip,I++. Without indirection,
 404 the STA ip,I++ instruction will write directly into the FORTH VM instruction
 405 stream. With indirection, the STA ip,I++ instruction can be used for directly
 406 updating byte/word variables whose pointers are stored directly in the FORTH
 407 VM instruction stream. (Although, the ability to create may be useful when
 408 compiling FORTH programs and for creating self-modifying FORTH programs, the
 409 STA ip,I++ instruction is expected to be prefixed with IND or ISZ under normal
 410 usage.)

411
 412 Finally, the ADD ip,I++ instruction allows constants (or relative offsets)
 413 located at the current IP to be added to the accumulator. Like LDA ip,I++ and
 414 STA ip,I++, the ADD ip,I++ supports the IND, SIz, and ISZ prefix instructions.

415
 416 Some consideration was given to directly supporting IP-relative conditional
 417 branches for the FORTH VM with the relative branch instructions of the
 418 M65C02A. Given that the ADD ip,I++ and LDA ip,I++ instructions can use the
 419 same microsequence, it was decided that directly supporting FORTH branches or
 420 jumps was too expensive in area and speed. IP-relative conditional FORTH
 421 branches can be implemented using the following instruction sequence:

```

422
423         [SIz] Bxx $1      ; [2[3]] test xx condition and branch if not true
424         ISZ DUP           ; [2] exchange A and IP (XAI)
425         SIz ADD ip,I++    ; [5] add IP-relative offset to A
426         ISZ DUP           ; [2] exchange A and IP (XAI)

```

427 \$1:

428
 429 The IP-relative conditional branch instruction sequence only requires 11[12]
 430 clock cycles, and IP-relative jumps require only 9 clock cycles. Conditional
 431 branches and unconditional jumps to absolute addresses rather than relative
 432 addresses can also be easily implemented. A conditional branch to an absolute
 433 address can be implemented as follows:

```

434
435         [SIz] Bxx $1      ; [2[3]] test xx condition and branch if not true
436         SIz LDA ip,I++    ; [5] load relative offset and autoincrement IP
437         IND DUP           ; [2] transfer A to IP (TAI)

```

438 \$1:

439
 440 Thus, a conditional branch to an absolute address requires 9[10] cycles, and
 441 the unconditional absolute jump only requires 7 clock cycles. Clearly, if the
 442 position independence of IP-relative branches and jumps is not required, then
 443 the absolute address branches and jumps provide greater performance.

444
 445 (Note: the M65C02A supports the eight 6502/65C02 branch instructions which
 446 perform true/false tests of the four ALU flags. When prefixed by the SIz, the
 447 eight branch instructions support additional tests of the ALU flags which
 448 support both signed and unsigned comparisons. The four signed conditional
 449 branches supported are: less than, less than or equal, greater than, and
 450 greater than or equal. The four unsigned conditional branches supported are:
 451 lower than, lower than or same, higher than, and higher than or same. These
 452 conditional branches are enabled by letting the 16-bit comparison instructions
 453 set the V flag.)

454
 455 The following table provides the instruction lengths (cycles) for the M65C02A-
 456 specific instructions which support the implementation of FORTH VMs:

	DTC	ITC	
459			
460	NXT	1(3)	; NEXT
461	ENT	1(5)	; ENTER/CALL/DOCOLON
462	PLI NXT	2(6)	; EXIT
463	--		
464	IND NXT	2(6)	
465	IND ENT	2(8)	
466	PLI IND NXT	3(9)	
467	--		
468	PLI	1(3)	; Pop IP
469	PHI	1(3)	; Push IP
470	INI	1(1)	; Increment IP
471	--		
472	IND PLI	2(4)	; PLW - Pop W
473	IND PHI	2(4)	; PHW - Push W
474	IND INI	2(2)	; INW - Increment W

```

475 --
476 LDA ip,I++      2(4)      ; Load byte from IP++ into A
477 SIZ LDA ip,I++   3(5)      ; Load word from IP++ into A
478 IND LDA ip,I++   3(7)      ; Load byte from IP++ indirect into A
479 ISZ LDA ip,I++   3(8)      ; Load word from IP++ indirect into A
480 --
481 STA ip,I++      2(4)      ; Store byte in A at IP++
482 SIZ STA ip,I++   3(5)      ; Store word in A at IP++
483 IND STA ip,I++   3(7)      ; Store byte in A at IP++ indirect
484 ISZ STA ip,I++   3(8)      ; Store word in A at IP++ indirect
485 --
486 ADD ip,I++      2(4)      ; Add byte from IP++ into A
487 SIZ ADD ip,I++   3(5)      ; Add word from IP++ into A
488 IND ADD ip,I++   3(7)      ; Add byte from IP++ indirect into A
489 ISZ ADD ip,I++   3(8)      ; Add word from IP++ indirect into A
490 --
491 IND DUP          2(2)      ; TAI - Transfer A to IP
492 SIZ DUP          2(2)      ; TIA - Transfer IP to A
493 ISZ DUP          2(2)      ; XAI - Exchange A and IP
494
495 The M65C02A implements a base-pointer relative addressing mode. This mode can
496 be used to directly access variables on the PS or the RS. Thus, the address
497 calculations required to access the PS and RS required when using a 6502/65C02
498 processor are not required with the M65C02A core. This addressing mode,
499 although also applicable to other HLLs, is expected to significantly improve
500 the performance of FORTH programs on the M65C02A core relative to the same
501 programs on 6502/65C02 FORTH VM implementations.
502
503 This concludes the FORTH VM trade study for the M65C02A soft-core processor.

```