

1 This document provides a description of the FORTH VM support being included in
2 the M65C02A soft-core processor. The M65C02A soft-core processor is an
3 extended version of the 65C02 microprocessor originally developed by
4 Western Design Center (WDC).
5

6 The purpose of this document is to document the design decisions made with
7 respect to custom instructions added to the basic M65C02A instruction set in
8 order to provide better support for a FORTH VM than a standard 65C02-
9 compatible processor. It is an objective of this effort to add the least
10 number of instructions (and dedicated hardware) to the M65C02A instruction set
11 as necessary to provide an efficient FORTH VM.
12

13 The instruction set of the M65C02A is already extended well beyond that of a
14 standard 6502/65C02 microprocessor. The M65C02A implements all of the standard
15 6502/65C02 instructions and addressing modes, but it also supports the four
16 bit-oriented instructions added by Rockwell and the WAI and STP instructions
17 included by WDC in its extended 6502/65C02 implementation: the W65C02S. In
18 addition, the bit-oriented RMBx/SMBx Rockwell instructions can be replaced
19 (through a microprogram change) by 16 M65C02A-specific instructions that
20 provide stack-relative and base-relative instructions.
21

22 The M65C02A provides 8 stack pointer-relative, and 8 post-indexed stack
23 pointer-relative indirect instructions matching the W65C816 processor's 16
24 stack-relative/stack-relative indirect instructions. In addition, the M65C02A
25 provides (by replacing the RMBx/SMBx bit-oriented Rockwell instructions) 8
26 base-relative instructions complementing the stack-relative instructions, and
27 an additional 8 stack-relative RMW instructions by replacing the RMBx/SMBx
28 instructions (through a microprogram change):
29

30 base-relative: ORA/AND/EOR/ADC/STA/LDA/CMP/SBC bp,B

31 stack-relative: ASL/ROL/LSR/ROR/TSB/TRB/INC/DEC sp,S
32

33 The M65C02A adds three W65C816 16-bit push instructions: PSH zp, PSH abs, and
34 PSH rel16. The M65C02A adds two 16-bit pull instructions: PUL zp, and PUL abs.
35 The M65C02A also adds signed/unsigned conditional/unconditional branch (8-bit
36 PC-relative) and jump (16-bit PC-relative) instructions, an instruction to
37 swap A with a 16-bit stack-relative location, and a block move instruction.
38 The M65C02A also includes support for multiple hardware-assisted stacks
39 (which can be greater than 256 bytes in size), Kernel/User mode with
40 independent system stack pointers (Sk/Su), an instruction to support co-
41 processors, and six prefix instructions that change the operand size, add
42 indirection, and override the destination registers of most instructions.
43

44 The IND prefix instruction provides the programmer with the ability to add
45 indirection to many instructions which do not include an indirect addressing
46 mode (including the bit-oriented Rockwell instructions). The SIZ and the
47 OAX/OAX prefix instructions provide the programmer the means to increase the
48 operand/ALU operation size, and to override the default destination register
49 of an instruction. The OSX prefix instruction allows the X register to be used
50 as an auxiliary hardware stack pointer. The effects of IND and SIZ are
51 combined in the ISZ prefix instruction.
52

53 The M65C02A soft-core processor is an implementation that provides a faithful
54 reimplement of the 6502/65C02 processor in a synthesizable manner.
55 Although it provides all of the registers and executes all of the standard
56 instructions, its implementation attempts to provide the best performance
57 possible. As a result, the M65C02A reimplement is not instruction cycle
58 length compatible with 6502/65C02 or W65C02S processors.
59

60 In all other respects it is a 6502/65C02 microprocessor. That means that, in
61 its 6502/65C02 compatibility mode, it appears to the programmer as having an
62 8-bit arithmetic accumulator (A), an 8-bit processor flags registers (P), two
63 8-bit index registers (X, Y), and a 8-bit system stack pointer (S). Except for

64 S, the general purpose accumulator (A) and the index registers (X, Y) do not
 65 provide any hardware support for the FORTH VM data or return stacks. In its
 66 extended mode, the M65C02A, has three modified 16-bit push-down stacks, one
 67 for each base register: A, X, Y. Each register stack is composed of three 16-
 68 bit registers arranged such that the top-of-stack is the register directly
 69 affected by the load/store and ALU instructions. Automatic pushing/popping of
 70 each register stack is not performed. Instead, explicit management of each
 71 register stack is left to the programmer. Three specific instructions are
 72 available for managing the register stacks: DUP (duplicate/push TOS), SWP
 73 (swap TOS and NOS), ROT (rotate the register stack). Thus, to load the TOS
 74 register and to push down the values currently in the stack, a DUP instruction
 75 must precede a load instruction. To simply duplicate the TOS and push down the
 76 stack, only a DUP instruction is needed. To store the TOS value to memory and
 77 pop the stack, a store instruction followed by a ROT instruction is needed.

78
 79 Although the extensions to the 6502/65C02 instruction set provided by the
 80 M65C02A are uniformly applicable to most instructions and registers, each
 81 register stack have some unique capabilities. For example, the A register
 82 stack supports byte swapping and bit reversal in the stack's TOS register.
 83 Similarly, the TOS of the X register stack can function as a third hardware-
 84 assisted stack pointer and as a automatically incremented or decremented
 85 memory pointer. Finally, the TOS of the Y register stack may function as an
 86 automatically incremented or decremented memory pointer for the block move
 87 instruction.

88
 89 The FORTH VM can generally be thought of as being constructed from a minimum
 90 of two stacks: (1) a parameter/data stack (PS), and (2) a return stack (RS).
 91 The PS is intended to hold all parameters/data used within a program/function,
 92 and the RS generally holds the return addresses of FORTH program words. In
 93 addition to holding FORTH program word return addresses, the RS is also used
 94 to hold loop addresses, and may be used to hold parameter/data addresses.

95
 96 These stacks are generally implemented within the memory of whatever
 97 microprocessor is implementing the FORTH VM. To support an efficient
 98 implementation of the FORTH VM, any specific FORTH implementation should
 99 provide hardware assisted stack pointers whenever possible.

100
 101 Performance degradations in a 6502/65C02 FORTH VM are generally due to the
 102 fact that all of the registers are 8-bits in length and 16 bits is the assumed
 103 operand and pointer size of the VM. Support is provided in the 6502/65C02
 104 instruction set architecture for multi-precision addition and subtraction, but
 105 any operations greater than 8 bits in length will entail several loads and
 106 stores. All of the additional steps necessary to implement a 16-bit or 32-bit
 107 FORTH VM operations reduces the performance a native 6502/65C02 FORTH VM can
 108 deliver.

109
 110 Brad Rodriguez wrote a series of articles on the development of FORTH VMs. The
 111 lead article of the series, "MOVING FORTH Part 1: Design Decisions in the
 112 Forth Kernel" (<http://www.bradrodriguez.com/papers/moving1.htm>), provides a
 113 good discussion of the design trades needed when implementing a FORTH VM on a
 114 microprocessor.

115
 116 Brad Rodriguez identifies the following registers as being the "classic" FORTH
 117 VM registers:

118		
119	W - Work Register	: general work register
120	IP - Interpreter Pointer	: address of the next FORTH word to execute
121	PSP - Parameter Stack Pointer	: points to the top of parameter/data stack
122	RSP - Return Stack Pointer	: points to the return address
123	UP - User Pointer	: points to User space of a multi-task FORTH
124	X - eXtra register	: temporary register for next address

125
 126 There are several generally accepted methods for implementing the FORTH VM.

127 The classic FORTH VM is implemented using a technique known as Indirect
 128 Threaded Code (ITC). The next most common approach is a technique known as
 129 Direct Threaded Code (DTC). A more recent approach is a technique known as
 130 Subroutine Threaded Code (STC). A final, less commonly used approach is a
 131 technique known as Token Threaded Code (TTC).
 132
 133 A TTC FORTH VM uses tokens that are smaller than the basic address pointer
 134 size to refer to FORTH words. The smaller size of the tokens allows a FORTH
 135 program to be compressed into a smaller image. The additional indirection
 136 required to locate the memory address of the FORTH word referenced by a
 137 particular token makes TTC FORTH VM implementations the slowest of the FORTH
 138 VM implementation techniques. The pre-indexed indirect addressing modes of the
 139 6502/65C02 can be used to implement a token threaded VM using the M65C02A.
 140 Thus, the basic indexed addressing modes of the M65C02A provide the necessary
 141 support to implement a TTC FORTH VM.
 142
 143 For an STC FORTH VM, each word is called using a native processor call
 144 instruction. Thus, there is no need for an IP register, and there is no inner
 145 interpreter. The FORTH program simply chains together the various FORTH words
 146 using native processor calls. There are two penalties for this simplicity: (1)
 147 subroutine calls are generally larger than simple address pointers; and (2)
 148 an STC FORTH requires pushing and popping the return stack on entry to and
 149 exit from each FORTH word. Thus, STC FORTH programs may be larger, and the
 150 additional push/pop operations performed by the native subroutine calling
 151 instructions may not deliver the performance improvements expected. Since an
 152 STC FORTH VM relies on the basic instructions of the processor, and the
 153 M65C02A provides those instructions, no additional instructions are needed in
 154 the M65C02A's instruction set to support an STC FORTH.
 155
 156 ITC FORTH and DTC FORTH VMs both require an inner interpreter. Thus, if the
 157 instruction set of a processor allows the implementation of the inner
 158 interpreter with a minimum number of instructions, then a FORTH VM on that
 159 processor would be "faster" than a FORTH VM on a processor without that
 160 support. This is the prime motivating factor for adding custom instructions to
 161 the M65C02A to support FORTH VMs.
 162
 163 A TTC/DTC/ITC FORTH VM is composed of two interpreters: (1) an outer
 164 interpreter, and (2) an inner interpreter. The outer interpreter is written in
 165 FORTH, i.e. it is composed of FORTH words. The inner interpreter, on the other
 166 hand, "executes" FORTH words. Therefore, a TTC/DTC/ITC FORTH VM spends the
 167 majority of its processing time in its inner interpreter. Thus, any decrease
 168 in the number of clock cycles required to "execute" a FORTH word will result
 169 in a clear increase in the performance of a FORTH program all other things
 170 being equal.
 171
 172 The basic structure of a 16-bit FORTH word is provided by the following C-like
 173 structure definition:
 174
 175 typedef struct {
 176 uint8_t Len;
 177 uint8_t [Max_Name_Len] Name;
 178 uint16_t *Link;
 179 uint16_t *Code_Fld;
 180 uint8_t [Code_Len] Param_Fld;
 181 } FORTH_word_t;
 182
 183 There are other forms, but the preceding structure defines all of the
 184 necessary components of the FORTH word, and succinctly conveys the required
 185 elements of a FORTH word.
 186
 187 The first three fields provide the "dictionary" header of FORTH words in a
 188 FORTH program. The first field defines the length of the name of the FORTH
 189 word. This is used to distinguish two or more FORTH words in a FORTH program

190 whose names share the same initial letters but which differ in length. The
191 second field defines the significant elements of the name of the FORTH word.
192 The total lengths of these two fields will determine the amount of memory that
193 is required just for the dictionary of a FORTH program. These two fields are
194 generally limited in size in order to conserve memory. If the immediate mode
195 of the FORTH compiler is not supported or included in the distribution of a
196 FORTH program, then the fields supporting the "dictionary" can be removed to
197 recover their memory for use by the application.
198
199 The fields, Code_Fld and Param_Fld, represent the "executable" portion of a
200 FORTH word. There are two types of FORTH words: (1) secondaries, and (2)
201 primitives. Secondaries are the predominant type of words in a FORTH program.
202 Their Param_Fld doesn't contain any native machine code. Instead, the
203 Param_Fld of FORTH secondaries is simply a list of pointers to the Code_Fld of
204 other FORTH words. Primitives are the FORTH words that perform the actual work
205 of any FORTH program. The Code_Fld of a primitive is a pointer/link to their
206 Param_Fld, which contains the machine code that performs the work the
207 primitive FORTH word is expected to provide.
208
209 In FORTH, the outer interpreter is used to perform immediate operations, and
210 to construct, define, or compile other FORTH words. As already stated above,
211 the FORTH VM's outer interpreter is generally composed of secondary FORTH
212 words. After all of the FORTH words associated with a FORTH program have been
213 compiled, the outer interpreter simply transfers control to the inner
214 interpreter to "execute" the top-most FORTH word of the program.
215
216 The FORTH VM's inner interpreter "executes" FORTH words. Since the outer
217 interpreter is mostly composed of secondary FORTH words, the inner interpreter
218 must move through each word until a FORTH primitive is found, and then
219 transfer temporary control to the machine code. The machine code of the
220 primitive FORTH word must return control to the inner interpreter once it
221 completes its task.
222
223 The inner interpreter "executes" the FORTH word that IP points to. It must
224 advance the IP through the Param_Fld of a secondary FORTH word, and jump to
225 the machine code pointed to by the Code_Fld of a FORTH primitive. The Code_Fld
226 of a secondary does not point to the Param_Fld of the word. Instead it points
227 to an inner interpreter function that "enters" the Param_Fld, i.e. performs
228 the FORTH equivalent of a subroutine call. The Code_Fld of a primitive does
229 point to the Param_Fld, and the inner interpreter simply jumps to the machine
230 code stored in the Param_Fld of the word.
231
232 Thus, there are three fundamental operations that the inner interpreter of a
233 DTC/ITC FORTH VM must perform:
234
235 (1) NEXT : fetch the FORTH word addressed by IP; advance IP.
236 (2) ENTER : push IP; load IP with the Code_Fld value; perform NEXT.
237
238 Each primitive FORTH word must transfer control back to the inner interpreter
239 so that the next FORTH word can be "executed". This action may be described as:
240
241 (3) EXIT : restore IP; perform NEXT.
242
243 These three fundamental operations are very similar to the operations that the
244 host processor performs in executing its machine code. NEXT corresponds
245 directly to the normal processor instruction fetch/execute cycle. ENTER
246 corresponds directly to a processor subroutine call and EXIT corresponds
247 directly to a processor subroutine return.
248
249 As previously discussed, FORTH uses two stacks: parameter/data stack and
250 return stack. Thus, ENTER and EXIT save and restore the IP to/from the return
251 stack, respectively. Most processors implement a single hardware stack into
252 which both return addresses and data are written. FORTH maintains strict

253 separation between the parameter/data stack and the return stack because it
 254 uses a stack-based arithmetic architecture. Mixing return addresses and
 255 parameters/data on a single stack would complicate the passing and processing
 256 of parameters.

257
 258 Table 6.3.1 in Koopman's "Stack Computers", provides a summary the relative
 259 frequency of the most frequently used FORTH words for several FORTH
 260 applications:

261	262 Name	FRAC	LIFE	MATH	COMPILE	AVE
263	CALL	11.16%	12.73%	12.59%	12.36%	12.21%
264	EXIT	11.07%	12.72%	12.55%	10.60%	11.74%
265	VARIABLE	7.63%	10.30%	2.26%	1.65%	5.46%
266	@	7.49%	2.05%	0.96%	11.09%	5.40%
267	0BRANCH	3.39%	6.38%	3.23%	6.11%	4.78%
268	LIT	3.94%	5.22%	4.92%	4.09%	4.54%
269	+	3.41%	10.45%	0.60%	2.26%	4.18%
270	SWAP	4.43%	2.99%	7.00%	1.17%	3.90%
271	R>	2.05%	0.00%	11.28%	2.23%	3.89%
272	>R	2.05%	0.00%	11.28%	2.16%	3.87%
273	CONSTANT	3.92%	3.50%	2.78%	4.50%	3.68%
274	DUP	4.08%	0.45%	1.88%	5.78%	3.05%
275	ROT	4.05%	0.00%	4.61%	0.48%	2.29%
276	USER	0.07%	0.00%	0.06%	8.59%	2.18%
277	C@	0.00%	7.52%	0.01%	0.36%	1.97%
278	I	0.58%	6.66%	0.01%	0.23%	1.87%
279	=	0.33%	4.48%	0.01%	1.87%	1.67%
280	AND	0.17%	3.12%	3.14%	0.04%	1.61%
281	BRANCH	1.61%	1.57%	0.72%	2.26%	1.54%
282	EXECUTE	0.14%	0.00%	0.02%	2.45%	0.65%

283
 284 In table above, CALL corresponds to ENTER. As can be seen, the remaining
 285 common FORTH words are a combination of parameter stack operations (DUP, ROT,
 286 SWAP), parameter stack loads (VARIABLE, LIT, CONSTANT, @, C@), parameter stack
 287 arithmetic and logic operations (+, =, AND), parameter stack branching and
 288 looping (0BRANCH, BRANCH, I), and parameter and return stack operations (R>,
 289 >R), and special operations (USER, EXECUTE).

290
 291 With the previous discussion and the FORTH word frequency data in the
 292 preceding table it is easy to assert that the M65C02A instruction set should
 293 contain custom instructions for at least NEXT, ENTER, and EXIT. The question
 294 then becomes to what extent should these operations be supported? In other
 295 words, should they be supported by a single instruction each and should they
 296 be supported for both ITC and DTC FORTH VMs?

297
 298 The following pseudo code defines the operations for these three operations in
 299 terms of the ITC and the DTC models:

300		ITC		DTC
301				
302		=====		=====
303	NEXT:	W <= (IP++) -- Ld *Code_Fld	; W <= (IP++) -- Ld *Code_Fld	
304		PC <= (W) -- Jump Indirect	; PC <= W -- Jump Direct	
305		=====		=====
306	ENTER:	(RSP--) <= IP -- Push IP on RS	; (RSP--) <= IP -- Push IP on RS	
307		IP <= W + 2 -- => Param_Fld	; IP <= W + 2 -- => Param_Fld	
308	;NEXT			
309		W <= (IP++) -- Ld *Code_Fld	; W <= (IP++) -- Ld *Code_Fld	
310		PC <= (W) -- Jump Indirect	; PC <= W -- Jump Direct	
311		=====		=====
312	EXIT:			
313		IP <= (++RSP) -- Pop IP frm RS	; IP <= (++RSP) -- Pop IP frm RS	
314	;NEXT			
315		W <= (IP++) -- Ld *Code_Fld	; W <= (IP++) -- Ld *Code_Fld	

```

316      PC      <= (W)      -- Jump Indirect      ; PC      <= W      -- Jump Direct
317 =====
318
319 Except for the indirection needed for ITC, there are several key takeaways
320 from the side-by-side comparison provided above of the NEXT, ENTER, and EXIT
321 operations. First, ENTER and EXIT are essentially the same for ITC and DTC
322 FORTH VMs. Second, both ENTER and EXIT terminate with the operations
323 implemented by NEXT. Also note that EXIT is simply an RS pop operation followed
324 by a DTC/ITC NEXT operation.
325
326 Finally, there are two important observations made by Rodriguez:
327
328 (1) if W is left pointing to the Code_Fld of the word being executed, the
329 Param_Fld of a FORTH word being ENTERed can be found using the value in W;
330
331 (2) providing a second stack pointer for the RS is important and will greatly
332 improve the performance of the inner interpreter.
333
334 The preceding analysis and discussions set the stage for the critical design
335 decisions for an M65C02A FORTH VM:
336
337 (1) mapping the FORTH VM registers onto the M65C02A registers;
338
339 (2) determining how to modify the M65C02A to support the FORTH VM inner
340 interpreter operations.
341
342 The IP register is strictly used as the instruction pointer of the inner
343 interpreter. It cannot be assigned to the target processor's program counter,
344 but it does operate as such for the inner interpreter. An easy means for
345 including IP is to place it in zero page memory, but this means that several
346 memory cycles will be needed to increment, push, pop, or otherwise manipulate
347 its value. A better solution is to add a 16-bit register within the core.
348 Alternatively, IP may be accessed with whatever custom instructions are added
349 to the M65C02A instruction set to support NEXT, ENTER, and EXIT. In addition
350 to load and store operations, support should be provided to increment the IP
351 by 2.
352
353 Similarly, the W register is used strictly as a pointer for indirect access to
354 a FORTH word by the inner interpreter. It is only loaded indirectly from IP.
355 Like the IP register, it can easily be implemented in zero page memory, but
356 this means that several memory cycles will be needed to increment, push, pop,
357 or otherwise manipulate its value. Like the IP, the best way for the M65C02A
358 to support W is to include it in the processor core itself. Also, for it to be
359 effectively utilized, support should be provided to increment the W register
360 by 2.
361
362 The PS is used more often than the RS. Thus, it makes more sense, from a speed
363 perspective, to use the native M65C02A stack for the PS. Using a pre-indexed
364 zero page location for the RSP will slow the push and pop operations
365 significantly. Therefore, a better solution would be to use one of the index
366 registers as the RSP, and place the RS anywhere in memory, including page 0.
367
368 As the 6502/65C02 index register for the less frequently used pre-indexed
369 (direct and indirect) addressing modes, the X index register is the natural
370 choice to provide a hardware-assisted RSP. The auxiliary stack pointer
371 capabilities of the TOS of the X register stack easily allow X to be used as
372 the RSP. In addition, two instructions will be added to support pushing and
373 pulling the IP from either stack, and a single cycle instruction will be added
374 to increment the IP by 1. By overloading the IND prefix instruction, it is
375 possible to use the dedicated IP push, pop, and increment instructions to
376 perform the same operations with the W register.
377
378 (Note: instead incurring a byte/cycle penalty by requiring that the OSX prefix

```

379 instruction be used before any FORTH VM instructions that use the RS, the OSX
380 prefix instruction's effects have been redefined to override of the default
381 stack pointer of any instruction that utilizes the stack. The default stack
382 for the FORTH VM's ENT, PHI, PLI, PHW, and PLW instructions has been defined
383 as the auxiliary stack provided by X. This change in the behavior of OSX means
384 that only when the PS is needed will these five FORTH VM instruction require
385 the OSX prefix instruction. It also saves 1 byte/cycle for every access to the
386 RS.)

387
388 Thus, the FORTH VM supported by the M65C02A will provide the following mapping
389 of the various FORTH VM registers:

390
391 IP - Internal dedicated 16-bit register
392 W - Internal dedicated 16-bit register
393 PSP - System Stack Pointer (S), allocated in memory (page 1 an option)
394 RSP - Auxiliary Stack Pointer (X), allocated in memory (page 0 an option)
395 UP - Memory (page 0 an option)
396 X - Not needed, M or MAR can provide any temporary storage required

397
398 The FORTH VM will be supported by the M65C02A using five single byte dedicated
399 instructions: NXT (NEXT), ENT (ENTER), PLI (Pull IP), PHI (Push IP), and INI
400 (Increment IP by 1). The inner interpreter is implemented by the NXT
401 instruction. These five instructions can be prefixed by IND. The NXT and ENT
402 instructions directly support a DTC FORTH VM. When prefixed by IND, they
403 support an ITC FORTH VM. The DTC EXIT will be implemented using the PLI NXT
404 instruction sequence, and the ITC EXIT will be implemented using the PLI IND
405 NXT instruction sequence. Access and control of the IP is provided by the PHI,
406 PLI, and INI instructions. (When these instructions are prefixed by IND,
407 access and control of W is provided: PHW, PLW, and INW, respectively.)

408
409 Loading constants/literals is a frequent operation in FORTH programs. Thus,
410 support for efficient loading of in-line constants/literals relative to the IP
411 is included in the M65C02A. The LDA ip,I++ instruction will load the byte
412 which follows the current IP into the accumulator and advances the IP by 1. If
413 this instruction is prefixed by SIZ, then the word following the current IP is
414 loaded into the accumulator and the IP is advanced by 2. If prefixed by IND,
415 the instruction becomes LDA (ip,I++), which uses the 16-bit word following the
416 current IP as a byte pointer. The IP is advanced by 2, and the byte pointed to
417 by the IP-relative pointer is loaded into the accumulator. If prefixed by ISZ,
418 the word following the current IP is used as a word pointer to load a 16-bit
419 value into the accumulator, while the IP is advanced by 2.

420
421 The LDA ip,I++ instruction is matched by the STA ip,I++. Without indirection,
422 the STA ip,I++ instruction will write directly into the FORTH VM instruction
423 stream. With IND/ISZ applied, the resulting STA (ip,I++) instruction can be
424 used for directly updating byte/word variables whose pointers are stored
425 directly in the FORTH VM instruction stream. (Although it may be useful when
426 compiling FORTH programs and for creating self-modifying FORTH programs, the
427 STA ip,I++ instruction is expected to be prefixed with IND or ISZ under normal
428 usage.)

429
430 Finally, the ADD ip,I++ instruction allows constants (or relative offsets)
431 located at the current IP to be added to the accumulator. Like LDA ip,I++ and
432 STA ip,I++, the ADD ip,I++ supports the IND, SIZ, and ISZ prefix instructions.
433 (Note: the ADD ip,I++ instruction does not add the the carry bit so CLC is not
434 required before the addition operation.)

435
436 Some consideration was given to directly supporting IP-relative conditional
437 branches for the FORTH VM with the relative branch instructions of the
438 M65C02A. Given that the ADD ip,I++ and LDA ip,I++ instructions can use the
439 same microsequence, it was decided that directly supporting FORTH branches or
440 jumps was too expensive in area and speed. IP-relative conditional FORTH
441 branches can be implemented using the following instruction sequence:

```

442
443         [SIZ] Bxx $1      ; [2[3]] test xx condition and branch if not true
444         ISZ DUP           ; [2] exchange A and IP (XAI)
445         SIZ ADD ip,I++    ; [5] add IP-relative offset to A
446         ISZ DUP           ; [2] exchange A and IP (XAI)
447     $1:
448
449 The IP-relative conditional branch instruction sequence only requires 11[12]
450 clock cycles, and IP-relative jumps require only 9 clock cycles. (Note: the
451 register stack manipulation instructions discussed above have been extended to
452 support exchanging the A top-of-stack register and IP. ISZ DUP exchanges ATOS
453 and IP (XAI), SIZ DUP transfers IP into A (TIA), and IND DUP transfers A into
454 IP (TAI).)
455
456 Conditional branches and unconditional jumps to absolute addresses rather than
457 relative addresses can also be easily implemented. A conditional branch to an
458 absolute address can be implemented as follows:
459
460         [SIZ] Bxx $1      ; [2[3]] test xx condition and branch if not true
461         SIZ LDA ip,I++    ; [5] load relative offset and autoincrement IP
462         IND DUP           ; [2] transfer A to IP (TAI)
463     $1:
464
465 Thus, a conditional branch to an absolute address requires 9[10] cycles, and
466 the unconditional absolute jump only requires 7 clock cycles. Clearly, if the
467 position independence of IP-relative branches and jumps is not required, then
468 the absolute address branches and jumps provide greater performance.
469
470 (Note: the M65C02A supports the eight 6502/65C02 branch instructions which
471 perform true/false tests of the four ALU flags. When prefixed by the SIZ
472 instruction, the eight branch instructions support additional tests of the ALU
473 flags which support both signed and unsigned comparisons. The four signed
474 conditional branches supported are: less than, less than or equal, greater
475 than, and greater than or equal. The four unsigned conditional branches
476 supported are: lower than, lower than or same, higher than, and higher than or
477 same. These conditional branches are enabled because the 16-bit comparison
478 instructions (SIZ/ISZ CMP/CPX/CPY) set the V flag.)
479
480 The following table provides the instruction lengths (cycles) for the M65C02A-
481 specific instructions which support the implementation of FORTH VMs:
482
483         DTC      ITC
484
485 NXT          1(3)          ; NEXT
486 ENT          1(5)          ; ENTER/CALL/DOCOLON
487 PLI NXT      2(6)          ; EXIT
488 --
489 IND NXT          2(6)
490 IND ENT          2(8)
491 PLI IND NXT      3(9)
492 --
493 PLI          1(3)          ; Pop IP
494 PHI          1(3)          ; Push IP
495 INI          1(1)          ; Increment IP
496 --
497 IND PLI        2(4)          ; PLW - Pop W
498 IND PHI        2(4)          ; PHW - Push W
499 IND INI        2(2)          ; INW - Increment W
500 --
501 LDA ip,I++      2(4)          ; Load byte from IP++ into A
502 SIZ LDA ip,I++  3(5)          ; Load word from IP++ into A
503 IND LDA ip,I++  3(7)          ; Load byte from IP++ indirect into A
504 ISZ LDA ip,I++  3(8)          ; Load word from IP++ indirect into A

```



```

505  --
506  STA ip,I++          2(4)          ; Store byte in A at IP++
507  SIZ STA ip,I++      3(5)          ; Store word in A at IP++
508  IND STA ip,I++      3(7)          ; Store byte in A at IP++ indirect
509  ISZ STA ip,I++      3(8)          ; Store word in A at IP++ indirect
510  --
511  ADD ip,I++          2(4)          ; Add byte from IP++ into A
512  SIZ ADD ip,I++      3(5)          ; Add word from IP++ into A
513  IND ADD ip,I++      3(7)          ; Add byte from IP++ indirect into A
514  ISZ ADD ip,I++      3(8)          ; Add word from IP++ indirect into A
515  --
516  IND DUP             2(2)          ; TAI - Transfer A to IP
517  SIZ DUP             2(2)          ; TIA - Transfer IP to A
518  ISZ DUP             2(2)          ; XAI - Exchange A and IP
519
520  The M65C02A implements a stack-relative addressing mode. This mode can be used
521  to directly access variables on the PS or the RS. Thus, the address
522  calculations required to access the PS and RS required when using a 6502/65C02
523  processor are not required with the M65C02A core. This addressing mode,
524  although also useful other HLLs, is expected to significantly improve the
525  performance of FORTH programs on the M65C02A core relative to the same
526  programs on 6502/65C02 FORTH VM implementations.
527
528  Finally, since the bit-oriented conditional branch instructions are not often
529  used, the M65C02A can be configured to replace the BBRx/BBSx Rockwell
530  instructions with a full complement of IP-relative with auto-increment
531  addressing instructions: AND/ORa/EOR/ADD/STA/LDA/CMP/SUB ip,I++, and
532  ASL/ROL/LSR/ROR/TSB/TRB/DEC/INC ip,I++.
533
534  This concludes the FORTH VM trade study for the M65C02A soft-core processor.

```